

In [1]:

```
!nvidia-smi
!nvcc --version

Sat Jan 10 16:19:54 2026
+-----
| NVIDIA-SMI 550.54.15      Driver Version: 550.54.15      CUDA Version: 12.4
|
|-----+-----+-----+
| GPU  Name           Persistence-M | Bus-Id          Disp.A  | Volatile Uncorr. EC
| C   |
| Fan  Temp     Perf            Pwr:Usage/Cap |          Memory-Usage | GPU-Util  Compute M.
|       |               |                               |                   | MIG
M.  |
|=====+=====+=====+
|=|
|  0  Tesla T4           Off  | 00000000:00:04.0 Off  |
0  |
| N/A   63C     P8          11W /  70W |        0MiB / 15360MiB |      0%     Default
t  |
|       |               |                               |                   |
/A  |
+-----+-----+-----+
|-----+-----+-----+
| Processes:
|   GPU  GI  CI          PID  Type  Process name          GPU Memori
y  |
|       ID  ID
|   ======+=====
|=|
|   No running processes found
|
+-----+
---+
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2024 NVIDIA Corporation
Built on Thu_Jun_6_02:18:23_PDT_2024
Cuda compilation tools, release 12.5, v12.5.82
Build cuda_12.5.r12.5/compiler.34385749_0
```

filter

cpu + gpu

In [2]:

```
# @title
%%file /tmp/t2v1.cu

#include <stdio.h>
#include <cuda_runtime.h>
#include <stdlib.h>
#include <chrono>
#include <math.h>

#define _DEBUGY 1
```

```

#ifndef _DEBUGY
#define CHECK_ERR(err, a) { if (err != cudaSuccess) { \
    printf("%s(%d): %s\n", __FILE__, __LINE__, a); \
} \
}
#else
#define CHECK_ERR(err, a) {}
#endif

// SoA: separate arrays for R, G, B
// Input: r_in[], g_in[], b_in[]
// Output: r_out[], g_out[], b_out[]

__global__ void blurKernelSoA(
    const float* __restrict__ r_in, //restrict hint to the compiler that this pointer is
the only way to access the data it points to during the scope of the function.
    const float* __restrict__ g_in,
    const float* __restrict__ b_in,
    float* __restrict__ r_out,
    float* __restrict__ g_out,
    float* __restrict__ b_out,
    int width, int height)
{
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;

    if (x >= width || y >= height) return;

    auto clamp = [] (int v, int lo, int hi) -> int {
        return (v < lo) ? lo : (v > hi ? hi : v);
    };

    float sum_r = 0.0f, sum_g = 0.0f, sum_b = 0.0f;
    int count = 0;

    for (int dy = -1; dy <= 1; dy++) {
        for (int dx = -1; dx <= 1; dx++) {
            int nx = clamp(x + dx, 0, width - 1);
            int ny = clamp(y + dy, 0, height - 1);
            int idx = ny * width + nx;

            sum_r += r_in[idx];
            sum_g += g_in[idx];
            sum_b += b_in[idx];
            count++;
        }
    }

    int out_idx = y * width + x;
    r_out[out_idx] = sum_r / count;
    g_out[out_idx] = sum_g / count;
    b_out[out_idx] = sum_b / count;
}

// CPU version (SoA)
void cpuBlurSoA(
    const float* r_in, const float* g_in, const float* b_in,
    float* r_out, float* g_out, float* b_out,
    int width, int height)
{
    auto clamp = [] (int v, int lo, int hi) -> int {
        return (v < lo) ? lo : (v > hi ? hi : v);
    };

    for (int y = 0; y < height; y++) {
        for (int x = 0; x < width; x++) {
            float sum_r = 0.0f, sum_g = 0.0f, sum_b = 0.0f;
            int count = 0;

            for (int dy = -1; dy <= 1; dy++) {
                for (int dx = -1; dx <= 1; dx++) {

```

```

        int nx = clamp(x + dx, 0, width - 1);
        int ny = clamp(y + dy, 0, height - 1);
        int idx = ny * width + nx;

        sum_r += r_in[idx];
        sum_g += g_in[idx];
        sum_b += b_in[idx];
        count++;
    }
}

int out_idx = y * width + x;
r_out[out_idx] = sum_r / count;
g_out[out_idx] = sum_g / count;
b_out[out_idx] = sum_b / count;
}
}

//=====
=====

int main() {
//const int WIDTH = 512;
//const int HEIGHT = 512;
//const int WIDTH = 10240;
//const int HEIGHT = 10240;
const int WIDTH = 1024;
const int HEIGHT = 1024;
const int NUM_PIXELS = WIDTH * HEIGHT;
size_t size = NUM_PIXELS * sizeof(float);

printf("SoA Blur: Image size %dx%d (%d pixels)\n", WIDTH, HEIGHT, NUM_PIXELS);

// Timers
cudaEvent_t t_start, t_stop;
cudaEventCreate(&t_start);
cudaEventCreate(&t_stop);

// Host memory (pinned)
float *h_r_in, *h_g_in, *h_b_in;
float *h_r_out_gpu, *h_g_out_gpu, *h_b_out_gpu;
cudaMallocHost(&h_r_in, size);
cudaMallocHost(&h_g_in, size);
cudaMallocHost(&h_b_in, size);
cudaMallocHost(&h_r_out_gpu, size);
cudaMallocHost(&h_g_out_gpu, size);
cudaMallocHost(&h_b_out_gpu, size);

float *h_r_out_cpu = (float*)malloc(size);
float *h_g_out_cpu = (float*)malloc(size);
float *h_b_out_cpu = (float*)malloc(size);

if (!h_r_in || !h_g_in || !h_b_in ||
    !h_r_out_gpu || !h_g_out_gpu || !h_b_out_gpu ||
    !h_r_out_cpu || !h_g_out_cpu || !h_b_out_cpu) {
    printf("Host alloc failed!\n");
    return -1;
}

// Device memory
float *d_r_in, *d_g_in, *d_b_in;
float *d_r_out, *d_g_out, *d_b_out;
cudaMalloc(&d_r_in, size);
cudaMalloc(&d_g_in, size);
cudaMalloc(&d_b_in, size);
cudaMalloc(&d_r_out, size);
cudaMalloc(&d_g_out, size);
cudaMalloc(&d_b_out, size);
CHECK_ERR(cudaGetLastError(), "cudaMalloc device");

// Initialize input

```

```

    srand(42);
    for (int i = 0; i < NUM_PIXELS; i++) {
        h_r_in[i] = (float)rand() / RAND_MAX * 255.0f;
        h_g_in[i] = (float)rand() / RAND_MAX * 255.0f;
        h_b_in[i] = (float)rand() / RAND_MAX * 255.0f;
    }

    // --- CPU blur ---
    auto tcpu_start = std::chrono::high_resolution_clock::now();
    cpuBlurSoA(h_r_in, h_g_in, h_b_in,
               h_r_out_cpu, h_g_out_cpu, h_b_out_cpu,
               WIDTH, HEIGHT);
    auto tcpu_stop = std::chrono::high_resolution_clock::now();
    float cpu_ms = std::chrono::duration_cast<std::chrono::microseconds>(tcpu_stop - tcpu_start).count() / 1000.0f;
    printf("CPU time: %.3f ms\n", cpu_ms);

    // --- GPU blur ---
    cudaMemcpy(d_r_in, h_r_in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_g_in, h_g_in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b_in, h_b_in, size, cudaMemcpyHostToDevice);
    CHECK_ERR(cudaGetLastError(), "H2D input");

    dim3 blockSize(16, 16);
    dim3 gridSize((WIDTH + blockSize.x - 1) / blockSize.x,
                  (HEIGHT + blockSize.y - 1) / blockSize.y);

    cudaEventRecord(t_start);
    blurKernelSoA<<<gridSize, blockSize>>>(
        d_r_in, d_g_in, d_b_in,
        d_r_out, d_g_out, d_b_out,
        WIDTH, HEIGHT);
    CHECK_ERR(cudaGetLastError(), "blurKernelSoA launch");

    cudaMemcpy(h_r_out_gpu, d_r_out, size, cudaMemcpyDeviceToHost);
    cudaMemcpy(h_g_out_gpu, d_g_out, size, cudaMemcpyDeviceToHost);
    cudaMemcpy(h_b_out_gpu, d_b_out, size, cudaMemcpyDeviceToHost);
    CHECK_ERR(cudaGetLastError(), "D2H output");

    cudaEventRecord(t_stop);
    cudaEventSynchronize(t_stop);
    float gpu_ms = 0;
    cudaEventElapsedTime(&gpu_ms, t_start, t_stop);
    printf("GPU time: %.3f ms\n", gpu_ms);

    // --- Validation ---
    const float tolerance = 1e-4f;
    int errors = 0;
    float max_err = 0.0f;
    const int sample_count = 1000;

    for (int i = 0; i < sample_count; i++) {
        int idx = rand() % NUM_PIXELS;
        float dr = fabsf(h_r_out_gpu[idx] - h_r_out_cpu[idx]);
        float dg = fabsf(h_g_out_gpu[idx] - h_g_out_cpu[idx]);
        float db = fabsf(h_b_out_gpu[idx] - h_b_out_cpu[idx]);
        float err = fmaxf(fmaxf(dr, dg), db);

        if (err > tolerance) {
            errors++;
            if (errors <= 3) {
                printf("Mismatch at %d: GPU(%f, %f, %f) CPU(%f, %f, %f) diff=%f\n",
                       idx,
                       h_r_out_gpu[idx], h_g_out_gpu[idx], h_b_out_gpu[idx],
                       h_r_out_cpu[idx], h_g_out_cpu[idx], h_b_out_cpu[idx],
                       err);
            }
        }
        if (err > max_err) max_err = err;
    }
}

```

```

printf("\nSoA Validation: %d errors in %d samples, max error = %.6f\n",
      errors, sample_count, max_err);

// Cleanup
cudaFreeHost(h_r_in); cudaFreeHost(h_g_in); cudaFreeHost(h_b_in);
cudaFreeHost(h_r_out_gpu); cudaFreeHost(h_g_out_gpu); cudaFreeHost(h_b_out_gpu);
free(h_r_out_cpu); free(h_g_out_cpu); free(h_b_out_cpu);
cudaFree(d_r_in); cudaFree(d_g_in); cudaFree(d_b_in);
cudaFree(d_r_out); cudaFree(d_g_out); cudaFree(d_b_out);
cudaEventDestroy(t_start);
cudaEventDestroy(t_stop);

return 0;
}

```

Writing /tmp/t2v1.cu

In [3]:

```

# @title
!nvcc -arch=sm_75 /tmp/t2v1.cu -o /tmp/t2v1

```

In [4]:

```

# @title
# run the executable with nvprof
!nvprof /tmp/t2v1

```

```

SoA Blur: Image size 1024x1024 (1048576 pixels)
==570== NVPYPROF is profiling process 570, command: /tmp/t2v1
CPU time: 111.295 ms
GPU time: 1.392 ms

```

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		46.81%	1.0299ms	3	343.29us	341.40us	346.43us	[CUDA memcpy HtoD]
		43.66%	960.52us	3	320.17us	319.61us	320.51us	[CUDA memcp y DtoH]
		9.53%	209.76us	1	209.76us	209.76us	209.76us	blurKernels
oA(float const *, float const *, float const *, float*, float*, float*, int, int)								
API calls:		84.94%	119.12ms	2	59.562ms	842ns	119.12ms	cudaEventCr eate
		8.43%	11.827ms	6	1.9712ms	1.7140ms	2.9816ms	cudaHostAll oc
		2.35%	3.2933ms	6	548.88us	332.57us	1.3846ms	cudaMemcpy
		2.33%	3.2682ms	6	544.70us	491.19us	695.34us	cudaFreeHos t
		0.82%	1.1478ms	6	191.30us	101.18us	216.31us	cudaFree
PCIBusId		0.57%	799.71us	1	799.71us	799.71us	799.71us	cuDeviceGet
		0.29%	406.28us	6	67.713us	60.594us	89.897us	cudaMalloc
ernel		0.13%	179.43us	1	179.43us	179.43us	179.43us	cudaLaunchK
		0.12%	161.88us	114	1.4190us	109ns	71.247us	cuDeviceGet
Attribute		0.01%	12.683us	1	12.683us	12.683us	12.683us	cuDeviceGet
Name		0.01%	10.522us	2	5.2610us	3.6270us	6.8950us	cudaEventRe cord
		0.00%	6.0220us	1	6.0220us	6.0220us	6.0220us	cudaEventSynch ronize
apsedTime		0.00%	3.6140us	1	3.6140us	3.6140us	3.6140us	cudaEventEl
		0.00%	2.6210us	2	1.3100us	478ns	2.1430us	cudaEventDe stroy
Error		0.00%	2.0920us	4	523ns	154ns	1.0500us	cudaGetLast
Count		0.00%	1.4150us	3	471ns	146ns	846ns	cuDeviceGet

	0.00%	923ns	1	923ns	923ns	923ns	cuDeviceTot
alMem	0.00%	806ns	2	403ns	175ns	631ns	cuDeviceGet
	0.00%	458ns	1	458ns	458ns	458ns	cuModuleGet
LoadingMode	0.00%	274ns	1	274ns	274ns	274ns	cuDeviceGet
Uuid							

In []:

Urmat teplo

cpu

In [5]:

```
# @title
%%file /tmp/t1v1.cu

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <string.h>
#include <algorithm>
#include <chrono>

// Параметры задачи
#define NX 256           // Количество узлов по x
#define NY 256           // Количество узлов по y
#define DX (1.0/(NX-1)) // Шаг по пространству
#define DY (1.0/(NY-1)) // Шаг по времени
#define C 1.0             // Коэффициент теплопроводности

// Автоматический расчет устойчивого шага по времени
#define DT (0.25 * fmin(DX*DX, DY*DY) / C) // Условие Куранта

double** allocate_2d_array(int rows, int cols) {
    double** arr = (double**)malloc(rows * sizeof(double*));
    for (int i = 0; i < rows; i++) {
        arr[i] = (double*)calloc(cols, sizeof(double)); // Используем calloc для инициализации нулями
    }
    return arr;
}

void free_2d_array(double** arr, int rows) {
    for (int i = 0; i < rows; i++) {
        free(arr[i]);
    }
    free(arr);
}

// Инициализация начальных и граничных условий
void initialize(double** u) {
    // Начальное условие: u(x,y,0) = 1.0
    for (int i = 0; i < NX; i++) {
        for (int j = 0; j < NY; j++) {
            u[i][j] = 1.0;
        }
    }

    // Границные условия: u = 2.0 на всех границах
    for (int i = 0; i < NX; i++) {
        u[i][0] = 2.0;
        u[i][NY-1] = 2.0;
    }
}
```

```

for (int j = 0; j < NY; j++) {
    u[0][j] = 2.0;
    u[NX-1][j] = 2.0;
}

// Один шаг по времени методом крест
void time_step(double** u, double** u_new) {
    double dx2 = DX * DX;
    double dy2 = DY * DY;

    // Копируем граничные условия
    for (int i = 0; i < NX; i++) {
        u_new[i][0] = u[i][0];
        u_new[i][NY-1] = u[i][NY-1];
    }
    for (int j = 0; j < NY; j++) {
        u_new[0][j] = u[0][j];
        u_new[NX-1][j] = u[NX-1][j];
    }

    // Вычисление новых значений во внутренних точках
    for (int i = 1; i < NX - 1; i++) {
        for (int j = 1; j < NY - 1; j++) {
            double laplacian = (u[i-1][j] - 2*u[i][j] + u[i+1][j]) / dx2 +
                               (u[i][j-1] - 2*u[i][j] + u[i][j+1]) / dy2;

            u_new[i][j] = u[i][j] + DT * C * laplacian;

            // Проверка на численную устойчивость
            if (!std::isfinite(u_new[i][j])) {
                printf("Numerical instability at (%d, %d)! Value = %f\n", i, j, u_new[i][j]);
                return;
            }
        }
    }
}

// Вычисление максимального изменения за шаг (для мониторинга)
double max_change(double** u, double** u_new) {
    double max_diff = 0.0;
    for (int i = 1; i < NX - 1; i++) {
        for (int j = 1; j < NY - 1; j++) {
            double diff = fabs(u_new[i][j] - u[i][j]);
            if (diff > max_diff) {
                max_diff = diff;
            }
        }
    }
    return max_diff;
}

void save_to_file(double** u, const char* filename, double t) {
    FILE* file = fopen(filename, "w");
    if (!file) {
        printf("Error opening file %s\n", filename);
        return;
    }

    fprintf(file, "# Time: %.6f\n", t);
    fprintf(file, "# X Y U\n");

    for (int i = 0; i < NX; i++) {
        for (int j = 0; j < NY; j++) {
            fprintf(file, "% .6f % .6f % .6f\n", i*DX, j*DY, u[i][j]);
        }
    }

    fclose(file);
    printf("Saved to %s\n", filename);
}

```

```

int main(int argc, char* argv[]) {
    double target_time = 0.1;
    if (target_time <= 0) {
        printf("Error: target_time must be positive\n");
        return 1;
    }

    printf("Solving heat equation until t = %.6f\n", target_time);
    printf("Grid: %dx%d, dx=%f, dy=%f\n", NX, NY, DX, DY);
    printf("Automatic time step: dt=%f (for stability)\n", DT);
    printf("Courant number: %.6f\n", C * DT * (1.0/(DX*DX) + 1.0/(DY*DY)));

    double** u_current = allocate_2d_array(NX, NY);
    double** u_next = allocate_2d_array(NX, NY);
    initialize(u_current);
    initialize(u_next);

    double t = 0.0;
    int step = 0;
    printf("Step 0: t = %.6f\n", t);
    auto start_computation = std::chrono::high_resolution_clock::now();

    // Основной цикл по времени
    while (t < target_time) {
        // Если осталось меньше чем DT, уменьшаем шаг
        double actual_dt = DT;
        if (t + DT > target_time) {
            actual_dt = target_time - t;
        }

        time_step(u_current, u_next);
        std::swap(u_current, u_next);

        t += actual_dt;
        step++;

        // Вывод прогресса каждые 1000 шагов
        if (step % 1000 == 0) {
            double max_diff = max_change(u_next, u_current);
            printf("Step %d: t = %.6f, max change = %.8f\n", step, t, max_diff);
        }

        // Аварийный выход при неустойчивости
        if (step > 1000000) { // Защита от бесконечного цикла
            printf("Too many steps, possible instability!\n");
            break;
        }
    }

    auto end_computation = std::chrono::high_resolution_clock::now();

    printf("Final: step %d, t = %.6f\n", step, t);

    save_to_file(u_current, "solution_final.txt", t);

    // Симметричный вывод среза для проверки (исправленный)
    printf("\nSymmetric slice at y = %d (middle):\n", NY/2);
    printf("Left part (from center to left boundary):\n");

    // Используем безопасные границы
    int center = NX / 2;
    int max_offset = center; // максимальный offset до левой границы

    for (int offset = 0; offset <= max_offset; offset += NX/16) {
        int left_x = center - offset;
        int right_x = center + offset;

        // Проверяем границы для безопасности
        if (left_x >= 0 && left_x < NX && right_x >= 0 && right_x < NX) {
            printf(" u[%3d] [%3d] = %8.6f | u[%3d] [%3d] = %8.6f\n",
                left_x, NY/2, u_current[left_x][NY/2],

```

```

        right_x, NY/2, u_current[right_x][NY/2]);
    }

// Проверка симметрии более детально (только для безопасных offset)
printf("\nDetailed symmetry check around center:\n");
for (int offset = 1; offset <= 5; offset++) {
    int left_x = center - offset;
    int right_x = center + offset;

    if (left_x >= 0 && right_x < NX) {
        double left_val = u_current[left_x][NY/2];
        double right_val = u_current[right_x][NY/2];
        double symmetry_error = fabs(left_val - right_val);
        printf(" Offset %d: u[%d]=%.6f, u[%d]=%.6f, error=%.8f\n",
               offset, left_x, left_val, right_x, right_val, symmetry_error);
    }
}

printf("\nBoundary check:\n");
printf("Top-left: u[0][%d] = %.6f (should be 2.0)\n", NY-1, u_current[0][NY-1]);
printf("Center: u[%d][%d] = %.6f\n", NX/2, NY/2, u_current[NX/2][NY/2]);
printf("Bottom-right: u[%d][0] = %.6f (should be 2.0)\n", NX-1, u_current[NX-1][0]);

auto computation_time = std::chrono::duration_cast<std::chrono::microseconds>(end_computation - start_computation);
printf("\n==== Performance Results ====\n");
printf("Computation time: %.3f ms\n", computation_time.count() / 1000.0);
printf("Time per step: %.3f microseconds\n", computation_time.count() / (double)step);
}

free_2d_array(u_current, NX);
free_2d_array(u_next, NX);

return 0;
}

```

Writing /tmp/t1v1.cu

In [6]:

```
# @title
!nvcc -arch=sm_75 /tmp/t1v1.cu -o /tmp/t1v1
# run the executable with nvprof
!nvprof /tmp/t1v1
```

```
Solving heat equation until t = 0.100000
Grid: 256x256, dx=0.003922, dy=0.003922
Automatic time step: dt=0.00000384 (for stability)
Courant number: 0.500000
Step 0: t = 0.000000
Step 1000: t = 0.003845, max change = 0.00036008
Step 2000: t = 0.007689, max change = 0.00018007
Step 3000: t = 0.011534, max change = 0.00012005
Step 4000: t = 0.015379, max change = 0.00009007
Step 5000: t = 0.019223, max change = 0.00007232
Step 6000: t = 0.023068, max change = 0.00006113
Step 7000: t = 0.026913, max change = 0.00005419
Step 8000: t = 0.030757, max change = 0.00005036
Step 9000: t = 0.034602, max change = 0.00004903
Step 10000: t = 0.038447, max change = 0.00004854
Step 11000: t = 0.042291, max change = 0.00004716
Step 12000: t = 0.046136, max change = 0.00004521
Step 13000: t = 0.049981, max change = 0.00004294
Step 14000: t = 0.053825, max change = 0.00004051
Step 15000: t = 0.057670, max change = 0.00003803
Step 16000: t = 0.061515, max change = 0.00003559
Step 17000: t = 0.065359, max change = 0.00003321
Step 18000: t = 0.069204, max change = 0.00003094
Step 19000: t = 0.073049, max change = 0.00002879
Step 20000: t = 0.076894, max change = 0.00002676
Step 21000: t = 0.080738, max change = 0.00002485
```

```
Step 22000: t = 0.084583, max change = 0.00002307
Step 23000: t = 0.088428, max change = 0.00002141
Step 24000: t = 0.092272, max change = 0.00001986
Step 25000: t = 0.096117, max change = 0.00001842
Step 26000: t = 0.099962, max change = 0.00001708
Final: step 26010, t = 0.100000
Saved to solution_final.txt
```

```
Symmetric slice at y = 128 (middle):
Left part (from center to left boundary):
u[128][128] = 1.774887 | u[128][128] = 1.774887
u[112][128] = 1.778971 | u[144][128] = 1.779514
u[ 96][128] = 1.791608 | u[160][128] = 1.792673
u[ 80][128] = 1.812313 | u[176][128] = 1.813859
u[ 64][128] = 1.840289 | u[192][128] = 1.842256
u[ 48][128] = 1.874457 | u[208][128] = 1.876769
u[ 32][128] = 1.913495 | u[224][128] = 1.916063
u[ 16][128] = 1.955891 | u[240][128] = 1.958615
```

```
Detailed symmetry check around center:
```

```
Offset 1: u[127]=1.774887, u[129]=1.774921, error=0.00003413
Offset 2: u[126]=1.774921, u[130]=1.774989, error=0.00006826
Offset 3: u[125]=1.774989, u[131]=1.775092, error=0.00010238
Offset 4: u[124]=1.775092, u[132]=1.775228, error=0.00013648
Offset 5: u[123]=1.775228, u[133]=1.775399, error=0.00017057
```

```
Boundary check:
```

```
Top-left: u[0][255] = 2.000000 (should be 2.0)
Center: u[128][128] = 1.774887
Bottom-right: u[255][0] = 2.000000 (should be 2.0)
```

```
==== Performance Results ====
Computation time: 30785.412 ms
Time per step: 1183.599 microseconds
===== Warning: No profile data collected.
```

In [7]:

```
# @title
import numpy as np
import matplotlib.pyplot as plt

def quick_plot():
    try:
        data = np.loadtxt('solution_final.txt', comments='#')
        x = data[:, 0]
        y = data[:, 1]
        u = data[:, 2]

        # Преобразование в сетку
        x_unique = np.unique(x)
        y_unique = np.unique(y)
        nx = len(x_unique)
        ny = len(y_unique)

        u_grid = u.reshape(nx, ny)
        X, Y = np.meshgrid(x_unique, y_unique, indexing='ij')

        plt.figure(figsize=(12, 4))

        plt.subplot(1, 2, 1)
        #plt.contourf(X, Y, u_grid, levels=50, cmap='cividis')
        plt.contourf(X, Y, u_grid, levels=50, cmap='summer')
        #plt.contourf(X, Y, u_grid, levels=50, cmap='Reds')
        plt.colorbar(label='Temperature')
        plt.xlabel('X')
        plt.ylabel('Y')
        plt.title('Temperature Distribution')
        plt.axis('equal')

        plt.subplot(1, 2, 2)
        middle_y = ny // 2
```

```

        plt.plot(x_unique, u_grid[:, middle_y], 'b-', linewidth=2, label=f'Slice at y={y_
unique[middle_y]:.3f}')
        plt.xlabel('X')
        plt.ylabel('u(x)')
        plt.title('Middle Slice')
        plt.grid(True, alpha=0.3)
        plt.legend()

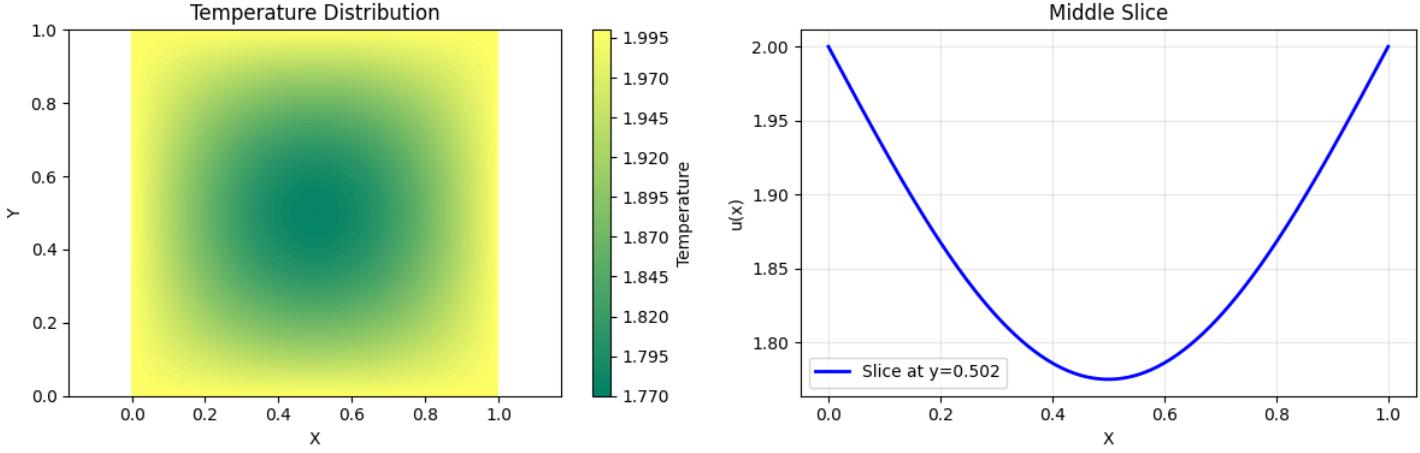
    plt.tight_layout()
    plt.show()

print(f"Min temperature: {u_grid.min():.4f}")
print(f"Max temperature: {u_grid.max():.4f}")
print(f"Center temperature: {u_grid[nx//2, ny//2]:.4f}")

except Exception as e:
    print(f"Error: {e}")

quick_plot()

```



Min temperature: 1.7749
 Max temperature: 2.0000
 Center temperature: 1.7749

In []:

gpu

In [8]:

```

# @title
%%file /tmp/t2v1.cu

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda_runtime.h>
#include <algorithm>
#include <chrono>

// Параметры задачи
#define NX 256
#define NY 256
#define DX (1.0f/(NX-1))
#define DY (1.0f/(NY-1))
#define C 1.0f
#define DT (0.25f * fminf(DX*DX, DY*DY) / C)

// Размер блока-----
#define BLOCK_SIZE 16

#define CHECK_CUDA_ERROR(err) \
if (err != cudaSuccess) { \

```

```

printf("CUDA error: %s at %s:%d\n", cudaGetErrorString(err), __FILE__, __LINE__);
exit(1); \

// CUDA kernel для одного шага по времени
__global__ void time_step_kernel(double* u, double* u_new, int width, int height,
                                 double dx2, double dy2, double dt_c) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= width || j >= height) return;

    int idx = j * width + i;

    // Границные условия
    if (i == 0 || i == width-1 || j == 0 || j == height-1) {
        u_new[idx] = u[idx];
        return;
    }

    // Вычисление лапласиана для внутренних точек
    double laplacian = (u[idx - 1] - 2*u[idx] + u[idx + 1]) / dx2 +
                       (u[idx - width] - 2*u[idx] + u[idx + width]) / dy2;

    u_new[idx] = u[idx] + dt_c * C * laplacian;
}

// CUDA kernel для инициализации
__global__ void initialize_kernel(double* u, int width, int height) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= width || j >= height) return;

    int idx = j * width + i;
    u[idx] = 1.0;

    if (i == 0 || i == width-1 || j == 0 || j == height-1) {
        u[idx] = 2.0;
    }
}

double compute_max_change_cpu(double* u, double* u_new, int width, int height) {
    double max_diff = 0.0;
    for (int j = 1; j < height-1; j++) {
        for (int i = 1; i < width-1; i++) {
            int idx = j * width + i;
            double diff = fabs(u_new[idx] - u[idx]);
            if (diff > max_diff) {
                max_diff = diff;
            }
        }
    }
    return max_diff;
}

double* allocate_gpu_memory(int size) {
    double* ptr;
    CHECK_CUDA_ERROR(cudaMalloc(&ptr, size * sizeof(double)));
    return ptr;
}

void free_gpu_memory(double* ptr) {
    CHECK_CUDA_ERROR(cudaFree(ptr));
}

void copy_between_cpu_gpu(double* dst, double* src, int size, bool to_gpu) {
    cudaMemsetKind kind = to_gpu ? cudaMemcpyHostToDevice : cudaMemcpyDeviceToHost;
    CHECK_CUDA_ERROR(cudaMemcpy(dst, src, size * sizeof(double), kind));
}

```

```

double** allocate_2d_array(int rows, int cols) {
    double** arr = (double**)malloc(rows * sizeof(double*));
    for (int i = 0; i < rows; i++) {
        arr[i] = (double*)calloc(cols, sizeof(double));
    }
    return arr;
}

void free_2d_array(double** arr, int rows) {
    for (int i = 0; i < rows; i++) {
        free(arr[i]);
    }
    free(arr);
}

void save_to_file(double* u, const char* filename, double t, int width, int height) {
    FILE* file = fopen(filename, "w");
    if (!file) {
        printf("Error opening file %s\n", filename);
        return;
    }

    fprintf(file, "# Time: %.6f\n", t);
    fprintf(file, "# X Y U\n");

    for (int i = 0; i < width; i++) {
        for (int j = 0; j < height; j++) {
            int idx = j * width + i;
            fprintf(file, "% .6f % .6f % .6f\n", i*DX, j*DY, u[idx]);
        }
    }

    fclose(file);
    printf("Saved to %s\n", filename);
}

int main() {
    double target_time = 0.1;

    printf("Solving heat equation with CUDA (SYMMETRIC OUTPUT) until t = %.6f\n", target_time);
    printf("Grid: %dx%d, dx=% .6f, dy=% .6f\n", NX, NY, DX, DY);
    printf("Time step: dt=% .8f\n", DT);
    printf("Block size: %dx%d\n", BLOCK_SIZE, BLOCK_SIZE);

    double* h_u = (double*)malloc(NX * NY * sizeof(double));
    double* h_u_new = (double*)malloc(NX * NY * sizeof(double));
    double* d_u = allocate_gpu_memory(NX * NY);
    double* d_u_new = allocate_gpu_memory(NX * NY);

    dim3 blockSize(BLOCK_SIZE, BLOCK_SIZE);
    dim3 gridSize((NX + BLOCK_SIZE - 1) / BLOCK_SIZE,
                  (NY + BLOCK_SIZE - 1) / BLOCK_SIZE);

    printf("Grid size: %dx%d blocks\n", gridSize.x, gridSize.y);

    for (int i = 0; i < NX * NY; i++) {
        h_u[i] = 1.0;
    }
    // Границные условия
    for (int i = 0; i < NX; i++) {
        h_u[i] = 2.0;
        h_u[(NY-1)*NX + i] = 2.0;
    }
    for (int j = 0; j < NY; j++) {
        h_u[j*NX] = 2.0;
        h_u[j*NX + (NX-1)] = 2.0;
    }

    copy_between_cpu_gpu(d_u, h_u, NX * NY, true);
    copy_between_cpu_gpu(d_u_new, h_u, NX * NY, true);
}

```

```

double t = 0.0;
int step = 0;

printf("Step 0: t = %.6f\n", t);

auto start_computation = std::chrono::high_resolution_clock::now();
while (t < target_time) {
    double actual_dt = DT;
    if (t + DT > target_time) {
        actual_dt = target_time - t;
    }

    time_step_kernel<<<gridSize, blockSize>>>(d_u, d_u_new, NX, NY, DX*DX, DY*DY, ac-
tual_dt);
    CHECK_CUDA_ERROR(cudaDeviceSynchronize());

    std::swap(d_u, d_u_new);
    t += actual_dt;
    step++;

    // Проверка максимального изменения каждые 1000 шагов
    if (step % 1000 == 0) {
        // Копируем оба массива на CPU для вычисления max_change
        copy_between_cpu_gru(h_u, d_u, NX * NY, false);
        copy_between_cpu_gru(h_u_new, d_u_new, NX * NY, false);

        double max_diff = compute_max_change_cpu(h_u, h_u_new, NX, NY);
        printf("Step %d: t = %.6f, max change = %.8f\n", step, t, max_diff);
    }

    if (step > 1000000) {
        printf("Too many steps, possible instability!\n");
        break;
    }
}

auto end_computation = std::chrono::high_resolution_clock::now();

printf("Final: step %d, t = %.6f\n", step, t);

// Копируем результат обратно на CPU
copy_between_cpu_gru(h_u, d_u, NX * NY, false);
save_to_file(h_u, "solution_final_cuda_symmetric.txt", t, NX, NY);

printf("\nSymmetric slice at y = %d (middle):\n", NY/2);
printf("Left part (from center to left boundary):\n");
for (int offset = 0; offset <= NX/2; offset += NX/16) {
    int left_idx = (NY/2) * NX + (NX/2 - offset);
    int right_idx = (NY/2) * NX + (NX/2 + offset);
    printf(" u[%3d][%3d] = %.6f | u[%3d][%3d] = %.6f\n",
           NX/2 - offset, NY/2, h_u[left_idx],
           NX/2 + offset, NY/2, h_u[right_idx]);
}

// Проверка симметрии более детально
printf("\nDetailed symmetry check around center:\n");
int center = NX / 2;
for (int offset = 1; offset <= 5; offset++) {
    int left_idx = (NY/2) * NX + (center - offset);
    int right_idx = (NY/2) * NX + (center + offset);
    double left_val = h_u[left_idx];
    double right_val = h_u[right_idx];
    double symmetry_error = fabs(left_val - right_val);
    printf(" Offset %d: u[%d]=%.6f, u[%d]=%.6f, error=%.8f\n",
           offset, center-offset, left_val, center+offset, right_val, symmetry_error
    );
}

printf("\nBoundary check:\n");
printf("Top-left corner: u[0][%d] = %.6f (should be 2.0)\n", NY-1, h_u[(NY-1)*NX
+ 0]);
printf("Center: u[%d][%d] = %.6f\n", NX/2, NY/2, h_u[(NY/2)*NX + NX/2]

```

```

);
printf("Bottom-right corner: u[%d][0] = %.6f (should be 2.0)\n", NX-1, h_u[0*NX + (N
X-1)]);

auto computation_time = std::chrono::duration_cast<std::chrono::microseconds>(end_com-
putation - start_computation);
printf("\n==== CUDA Performance Results ====\n");
printf("Computation time: %.3f ms\n", computation_time.count() / 1000.0);
printf("Time per step: %.3f microseconds\n", computation_time.count() / (double)step
);

free(h_u);
free(h_u_new);
free_gpu_memory(d_u);
free_gpu_memory(d_u_new);

return 0;
}

```

Overwriting /tmp/t2v1.cu

In [9]:

```

# @title
!nvcc -arch=sm_75 /tmp/t2v1.cu -o /tmp/t2v1
# run the executable with nvprof
!nvprof /tmp/t2v1

```

```

Solving heat equation with CUDA (SYMMETRIC OUTPUT) until t = 0.100000
Grid: 256x256, dx=0.003922, dy=0.003922
Time step: dt=0.00000384
Block size: 16x16
==1026== NVPROF is profiling process 1026, command: /tmp/t2v1
Grid size: 16x16 blocks
Step 0: t = 0.000000
Step 1000: t = 0.003845, max change = 0.00036008
Step 2000: t = 0.007689, max change = 0.00018007
Step 3000: t = 0.011534, max change = 0.00012005
Step 4000: t = 0.015379, max change = 0.00009007
Step 5000: t = 0.019223, max change = 0.00007232
Step 6000: t = 0.023068, max change = 0.00006113
Step 7000: t = 0.026913, max change = 0.00005419
Step 8000: t = 0.030757, max change = 0.00005036
Step 9000: t = 0.034602, max change = 0.00004903
Step 10000: t = 0.038447, max change = 0.00004854
Step 11000: t = 0.042291, max change = 0.00004716
Step 12000: t = 0.046136, max change = 0.00004521
Step 13000: t = 0.049981, max change = 0.00004294
Step 14000: t = 0.053825, max change = 0.00004051
Step 15000: t = 0.057670, max change = 0.00003803
Step 16000: t = 0.061515, max change = 0.00003559
Step 17000: t = 0.065359, max change = 0.00003321
Step 18000: t = 0.069204, max change = 0.00003094
Step 19000: t = 0.073049, max change = 0.00002879
Step 20000: t = 0.076894, max change = 0.00002676
Step 21000: t = 0.080738, max change = 0.00002485
Step 22000: t = 0.084583, max change = 0.00002307
Step 23000: t = 0.088428, max change = 0.00002141
Step 24000: t = 0.092272, max change = 0.00001986
Step 25000: t = 0.096117, max change = 0.00001842
Step 26000: t = 0.099962, max change = 0.00001708
Final: step 26010, t = 0.100000
Saved to solution_final_cuda_symmetric.txt

```

```

Symmetric slice at y = 128 (middle):
Left part (from center to left boundary):
u[128][128] = 1.774887 | u[128][128] = 1.774887
u[112][128] = 1.778971 | u[144][128] = 1.779514
u[ 96][128] = 1.791608 | u[160][128] = 1.792673
u[ 80][128] = 1.812313 | u[176][128] = 1.813859
u[ 64][128] = 1.840289 | u[192][128] = 1.842256
u[ 48][128] = 1.874457 | u[208][128] = 1.876769

```

```
u[ 32][128] = 1.913495 | u[224][128] = 1.916063
u[ 16][128] = 1.955891 | u[240][128] = 1.958615
u[  0][128] = 2.000000 | u[256][128] = 2.000000
```

Detailed symmetry check around center:

```
Offset 1: u[127]=1.774887, u[129]=1.774921, error=0.00003413
Offset 2: u[126]=1.774921, u[130]=1.774989, error=0.00006826
Offset 3: u[125]=1.774989, u[131]=1.775092, error=0.00010238
Offset 4: u[124]=1.775092, u[132]=1.775228, error=0.00013648
Offset 5: u[123]=1.775228, u[133]=1.775399, error=0.00017057
```

Boundary check:

```
Top-left corner: u[0][255] = 2.000000 (should be 2.0)
Center: u[128][128] = 1.774887
Bottom-right corner: u[255][0] = 2.000000 (should be 2.0)
```

==== CUDA Performance Results ===

Computation time: 856.832 ms

Time per step: 32.942 microseconds

==1026== Profiling application: /tmp/t2v1

==1026== Profiling result:

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	99.55%	575.31ms		26010	22.118us	15.679us	84.030us	time_step_kernel(double*, double*, int, int, double, double)
	0.44%	2.5376ms		53	47.878us	42.239us	53.790us	[CUDA memcpy DtoH]
	0.02%	90.590us		2	45.295us	44.991us	45.599us	[CUDA memcpy HtoD]
API calls:	72.90%	670.87ms		26010	25.792us	3.0620us	1.8114ms	cudaDeviceSynchronize
	16.25%	149.51ms		26010	5.7480us	3.2570us	585.96us	cudaLaunchKernel
	9.68%	89.041ms		2	44.520ms	3.2050us	89.037ms	cudaMalloc
	1.13%	10.422ms		55	189.50us	135.05us	492.94us	cudaMemcpy
	0.02%	193.86us		2	96.930us	29.736us	164.12us	cudaFree
	0.02%	156.67us		114	1.3740us	104ns	66.910us	cuDeviceGetAttribute
Attribute Name	0.00%	14.630us		1	14.630us	14.630us	14.630us	cuDeviceGet
PCIBusId	0.00%	5.5570us		1	5.5570us	5.5570us	5.5570us	cuDeviceGet
Count	0.00%	2.3800us		3	793ns	132ns	1.9690us	cuDeviceGet
	0.00%	1.1110us		2	555ns	203ns	908ns	cuDeviceGet
allMem	0.00%	509ns		1	509ns	509ns	509ns	cuDeviceTotal
LoadingMode	0.00%	423ns		1	423ns	423ns	423ns	cuModuleGet
Uuid	0.00%	239ns		1	239ns	239ns	239ns	cuDeviceGet

In [10]:

```
# @title
import numpy as np
import matplotlib.pyplot as plt

def quick_plot1():
    try:
        data = np.loadtxt('solution_final_cuda_symmetric.txt', comments='#')
        x = data[:, 0]
        y = data[:, 1]
        u = data[:, 2]

        # Преобразование в сетку
        x_unique = np.unique(x)
        y_unique = np.unique(y)
        nx = len(x_unique)
        ny = len(y_unique)

        u_grid = u.reshape(nx, ny)
        X, Y = np.meshgrid(x_unique, y_unique, indexing='ij')
```

```

plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
# plt.contourf(X, Y, u_grid, levels=50, cmap='cividis')
plt.contourf(X, Y, u_grid, levels=50, cmap='summer')
# plt.contourf(X, Y, u_grid, levels=50, cmap='Reds')
plt.colorbar(label='Temperature')
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Temperature Distribution')
plt.axis('equal')

plt.subplot(1, 2, 2)
middle_y = ny // 2
plt.plot(x_unique, u_grid[:, middle_y], 'b-', linewidth=2, label=f'Slice at y={y_unique[middle_y]:.3f}')
plt.xlabel('X')
plt.ylabel('u(x)')
plt.title('Middle Slice')
plt.grid(True, alpha=0.3)
plt.legend()

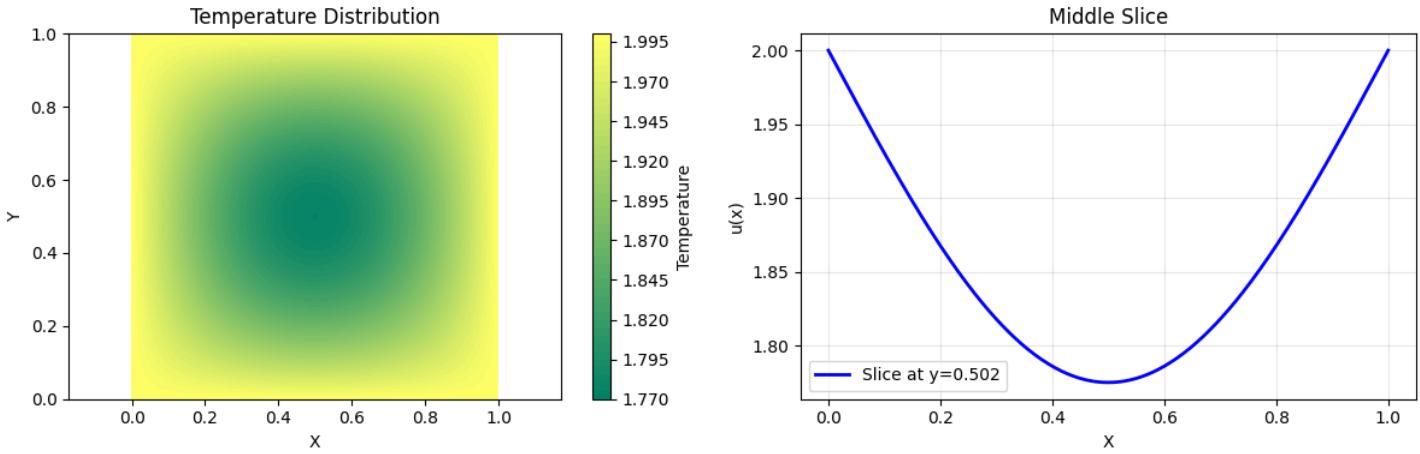
plt.tight_layout()
plt.show()

print(f"Min temperature: {u_grid.min():.4f}")
print(f"Max temperature: {u_grid.max():.4f}")
print(f"Center temperature: {u_grid[nx//2, ny//2]:.4f}")

except Exception as e:
    print(f"Error: {e}")

quick_plot1()

```



```

Min temperature: 1.7749
Max temperature: 2.0000
Center temperature: 1.7749

```

In []:

urmat laplac

cpu

In [11]:

```

# @title
%%file /tmp/tlvl.cu

#include <stdio.h>

```

```

#include <stdlib.h>
#include <math.h>
#include <string.h>
#include <algorithm>
#include <chrono>

// Параметры задачи
#define NX 256 // Количество узлов по x
#define NY 256 // Количество узлов по y
#define NX 128 // Количество узлов по x
#define NY 128 // Количество узлов по y
#define DX (1.0/(NX-1)) // Шаг по пространству
#define DY (1.0/(NY-1))
#define MAX_ITER 50000 // Максимальное количество итераций
#define TOLERANCE 1e-5 // Точность решения

double** allocate_2d_array(int rows, int cols) {
    double** arr = (double**)malloc(rows * sizeof(double*));
    for (int i = 0; i < rows; i++) {
        arr[i] = (double*)calloc(cols, sizeof(double));
    }
    return arr;
}

void free_2d_array(double** arr, int rows) {
    for (int i = 0; i < rows; i++) {
        free(arr[i]);
    }
    free(arr);
}

// Инициализация начальных и граничных условий
void initialize(double** u) {
    // Внутренняя область инициализируется нулями
    for (int i = 0; i < NX; i++) {
        for (int j = 0; j < NY; j++) {
            u[i][j] = 0.0;
        }
    }

    // Граничные условия из новой задачи
    for (int i = 0; i < NX; i++) {
        u[i][0] = exp(1.0 - i*DX); // u(x,0) = e^(1-x)
        u[i][NY-1] = 1.0; // u(x,1) = 1.0
    }

    for (int j = 0; j < NY; j++) {
        u[0][j] = exp(1.0 - j*DY); // u(0,y) = e^(1-y)
        u[NX-1][j] = 1.0; // u(1,y) = 1.0
    }
}

// Простой метод Якоби для уравнения Лапласа
void jacobi_step(double** u, double** u_new) {
    double dx2 = DX * DX;
    double dy2 = DY * DY;
    double factor = 0.5 / (1.0/dx2 + 1.0/dy2);

    // Копируем граничные условия
    for (int i = 0; i < NX; i++) {
        u_new[i][0] = u[i][0];
        u_new[i][NY-1] = u[i][NY-1];
    }
    for (int j = 0; j < NY; j++) {
        u_new[0][j] = u[0][j];
        u_new[NX-1][j] = u[NX-1][j];
    }

    // Вычисление новых значений во внутренних точках
    for (int i = 1; i < NX - 1; i++) {
        for (int j = 1; j < NY - 1; j++) {
            u_new[i][j] = factor * (

```

```

        (u[i-1][j] + u[i+1][j]) / dx2 +
        (u[i][j-1] + u[i][j+1]) / dy2
    );
}
}

// Вычисление максимального изменения за шаг
double max_change(double** u, double** u_new) {
    double max_diff = 0.0;
    for (int i = 1; i < NX - 1; i++) {
        for (int j = 1; j < NY - 1; j++) {
            double diff = fabs(u_new[i][j] - u[i][j]);
            if (diff > max_diff) {
                max_diff = diff;
            }
        }
    }
    return max_diff;
}

// Вычисление невязки (норма C) //-----
double compute_residual(double** u) {
    double max_residual = 0.0;
    double dx2 = DX * DX;
    double dy2 = DY * DY;

    for (int i = 1; i < NX - 1; i++) {
        for (int j = 1; j < NY - 1; j++) {
            double residual = fabs(
                (u[i-1][j] - 2*u[i][j] + u[i+1][j]) / dx2 +
                (u[i][j-1] - 2*u[i][j] + u[i][j+1]) / dy2
            );
            if (residual > max_residual) {
                max_residual = residual;
            }
        }
    }
    return max_residual;
}

// Вычисление нормы L1
double compute_l1_norm(double** u) {
    double sum = 0.0;
    for (int i = 0; i < NX; i++) {
        for (int j = 0; j < NY; j++) {
            sum += fabs(u[i][j]);
        }
    }
    return sum * DX * DY;
}

void save_to_file(double** u, const char* filename) {
    FILE* file = fopen(filename, "w");
    if (!file) {
        printf("Error opening file %s\n", filename);
        return;
    }

    fprintf(file, "# X Y U\n");
    for (int i = 0; i < NX; i++) {
        for (int j = 0; j < NY; j++) {
            fprintf(file, "% .6f % .6f % .6f\n", i*DX, j*DY, u[i][j]);
        }
    }

    fclose(file);
    printf("Saved to %s\n", filename);
}

int main(int argc, char* argv[]) {

```

```

printf("Solving Laplace equation: d2u/dx2 + d2u/dy2 = 0\n");
printf("Grid: %dx%d, dx=% .6f, dy=% .6f\n", NX, NY, DX, DY);
printf("Tolerance: %.2e, Max iterations: %d\n", TOLERANCE, MAX_ITER);

double** u_current = allocate_2d_array(NX, NY);
double** u_next = allocate_2d_array(NX, NY);
initialize(u_current);
initialize(u_next);

auto start_computation = std::chrono::high_resolution_clock::now();

// Основной итерационный цикл
int iter = 0;
double max_diff = 0.0;
double residual = 0.0;

printf("\nStarting iterations...\n");
printf("Iter Residual (C-norm)      Max-Change      L1-Norm\n");
printf("-----\n");

for (iter = 1; iter <= MAX_ITER; iter++) {
    jacobi_step(u_current, u_next);

    max_diff = max_change(u_current, u_next);
    residual = compute_residual(u_next);
    double l1_norm = compute_l1_norm(u_next);

    std::swap(u_current, u_next);

    // Вывод прогресса каждые 1000 итераций
    if (iter % 1000 == 0) {
        printf("%5d %12.6e %12.6e %12.6e\n",
               iter, residual, max_diff, l1_norm);
    }

    // Критерий остановки
    if (residual < TOLERANCE && max_diff < TOLERANCE) {
        printf("%5d %12.6e %12.6e %12.6e\n",
               iter, residual, max_diff, l1_norm);
        printf("Converged!\n");
        break;
    }
}

auto end_computation = std::chrono::high_resolution_clock::now();

if (iter > MAX_ITER) {
    printf("Reached maximum iterations without convergence\n");
    printf("Current residual: %.6e (target: %.1e)\n", residual, TOLERANCE);
}

printf("\n==== Final Results ====\n");
printf("Iterations: %d\n", iter);
printf("Final residual (C-norm): %.6e\n", residual);
printf("Final max change: %.6e\n", max_diff);
printf("Final L1 norm: %.6e\n", compute_l1_norm(u_current));

save_to_file(u_current, "solution_laplace.txt");

// Анализ решения
printf("\n==== Solution Analysis ====\n");

// Проверка граничных условий
printf("Boundary conditions check:\n");
printf("u(0,0) = %.6f (should be e = %.6f)\n", u_current[0][0], exp(1.0));
printf("u(1,0) = %.6f (should be 1.0)\n", u_current[NX-1][0]);
printf("u(0,1) = %.6f (should be 1.0)\n", u_current[0][NY-1]);
printf("u(1,1) = %.6f (should be 1.0)\n", u_current[NX-1][NY-1]);

// Проверка внутренних точек
printf("\nInternal points:\n");
printf("u(0.5, 0.5) = %.6f\n", u_current[NX/2][NY/2]);

```

```

printf("u(0.25, 0.25) = %.6f\n", u_current[NX/4][NY/4]);
printf("u(0.75, 0.75) = %.6f\n", u_current[3*NX/4][3*NY/4]);

auto computation_time = std::chrono::duration_cast<std::chrono::microseconds>(end_computation - start_computation);
printf("\n==== Performance Results ====\n");
printf("Total computation time: %.3f ms\n", computation_time.count() / 1000.0);
printf("Time per iteration: %.3f microseconds\n", computation_time.count() / (double)iter);
printf("Iterations per second: %.0f\n", iter / (computation_time.count() / 1000000.0));
);

free_2d_array(u_current, NX);
free_2d_array(u_next, NX);

return 0;
}

```

Overwriting /tmp/t1v1.cu

In [12]:

```

# @title
!nvcc -arch=sm_75 /tmp/t1v1.cu -o /tmp/t1v1
# run the executable with nvprof
!nvprof /tmp/t1v1

```

Solving Laplace equation: $d^2u/dx^2 + d^2u/dy^2 = 0$
Grid: 128x128, $dx=0.007874$, $dy=0.007874$
Tolerance: $1.00e-05$, Max iterations: 50000

Starting iterations...

Iter	Residual(C-norm)	Max-Change	L1-Norm
1000	4.825725e+01	7.485691e-04	6.747362e-01
2000	2.282065e+01	3.537575e-04	8.815974e-01
3000	1.647062e+01	2.553533e-04	1.015109e+00
4000	1.249145e+01	1.936543e-04	1.109617e+00
5000	9.287306e+00	1.439822e-04	1.178395e+00
6000	6.855223e+00	1.062780e-04	1.228867e+00
7000	5.049931e+00	7.829375e-05	1.265995e+00
8000	3.719094e+00	5.765806e-05	1.293328e+00
9000	2.738657e+00	4.245810e-05	1.313454e+00
10000	2.016613e+00	3.126406e-05	1.328275e+00
11000	1.484947e+00	2.302152e-05	1.339188e+00
12000	1.093468e+00	1.695284e-05	1.347225e+00
13000	8.052044e-01	1.248398e-05	1.353143e+00
14000	5.929389e-01	9.193138e-06	1.357501e+00
15000	4.366326e-01	6.769775e-06	1.360711e+00
16000	3.215317e-01	4.985224e-06	1.363074e+00
17000	2.367732e-01	3.671091e-06	1.364814e+00
18000	1.743580e-01	2.703370e-06	1.366096e+00
19000	1.283961e-01	1.990746e-06	1.367039e+00
20000	9.455003e-02	1.465974e-06	1.367734e+00
21000	6.962606e-02	1.079535e-06	1.368246e+00
22000	5.127221e-02	7.949636e-07	1.368623e+00
23000	3.775656e-02	5.854066e-07	1.368901e+00
24000	2.780371e-02	4.310901e-07	1.369105e+00
25000	2.047449e-02	3.174523e-07	1.369255e+00
26000	1.507730e-02	2.337701e-07	1.369366e+00
27000	1.110284e-02	1.721470e-07	1.369448e+00
28000	8.176065e-03	1.267680e-07	1.369508e+00
29000	6.020808e-03	9.335126e-08	1.369552e+00
30000	4.433688e-03	6.874334e-08	1.369585e+00
31000	3.264943e-03	5.062220e-08	1.369609e+00
32000	2.404285e-03	3.727790e-08	1.369626e+00
33000	1.770502e-03	2.745123e-08	1.369639e+00
34000	1.303788e-03	2.021493e-08	1.369649e+00
35000	9.601018e-04	1.488616e-08	1.369656e+00
36000	7.070135e-04	1.096208e-08	1.369661e+00
37000	5.206407e-04	8.072417e-09	1.369665e+00
38000	3.833969e-04	5.944482e-09	1.369668e+00

```
39000 2.823313e-04 4.377483e-09 1.369670e+00
40000 2.079071e-04 3.223553e-09 1.369672e+00
41000 1.531016e-04 2.373806e-09 1.369673e+00
42000 1.127432e-04 1.748057e-09 1.369674e+00
43000 8.302344e-05 1.287260e-09 1.369674e+00
44000 6.113799e-05 9.479306e-10 1.369675e+00
45000 4.502170e-05 6.980512e-10 1.369675e+00
46000 3.315375e-05 5.140413e-10 1.369675e+00
47000 2.441420e-05 3.785368e-10 1.369675e+00
48000 1.797849e-05 2.787524e-10 1.369675e+00
49000 1.323928e-05 2.052720e-10 1.369676e+00
49918 9.997030e-06 1.550018e-10 1.369676e+00
```

Converged!

==== Final Results ===

```
Iterations: 49918
Final residual (C-norm): 9.997030e-06
Final max change: 1.550018e-10
Final L1 norm: 1.369676e+00
Saved to solution_laplace.txt
```

==== Solution Analysis ===

Boundary conditions check:

```
u(0,0) = 2.718282 (should be e = 2.718282)
u(1,0) = 1.000000 (should be 1.0)
u(0,1) = 1.000000 (should be 1.0)
u(1,1) = 1.000000 (should be 1.0)
```

Internal points:

```
u(0.5, 0.5) = 1.337383
u(0.25, 0.25) = 1.815447
u(0.75, 0.75) = 1.080496
```

==== Performance Results ===

```
Total computation time: 20807.800 ms
Time per iteration: 416.840 microseconds
Iterations per second: 2399
===== Warning: No profile data collected.
```

In [13]:

```
# @title
import numpy as np
import matplotlib.pyplot as plt

def quick_plot():
    try:
        data = np.loadtxt('solution_laplace.txt', comments='#')
        x = data[:, 0]
        y = data[:, 1]
        u = data[:, 2]

        # Преобразование в сетку
        x_unique = np.unique(x)
        y_unique = np.unique(y)
        nx = len(x_unique)
        ny = len(y_unique)

        u_grid = u.reshape(nx, ny)
        X, Y = np.meshgrid(x_unique, y_unique, indexing='ij')

        plt.figure(figsize=(12, 4))

        plt.subplot(1, 2, 1)
        #plt.contourf(X, Y, u_grid, levels=50, cmap='cividis')
        plt.contourf(X, Y, u_grid, levels=50, cmap='summer')
        #plt.contourf(X, Y, u_grid, levels=50, cmap='Reds')
        plt.colorbar(label='Temperature')
        plt.xlabel('X')
        plt.ylabel('Y')
        plt.title('Temperature Distribution')
        plt.axis('equal')
```

```

plt.subplot(1, 2, 2)
middle_y = ny // 2
plt.plot(x_unique, u_grid[:, middle_y], 'b-', linewidth=2, label=f'Slice at y={y_
unique[middle_y]:.3f}')
plt.xlabel('X')
plt.ylabel('u(x)')
plt.title('Middle Slice')
plt.grid(True, alpha=0.3)
plt.legend()

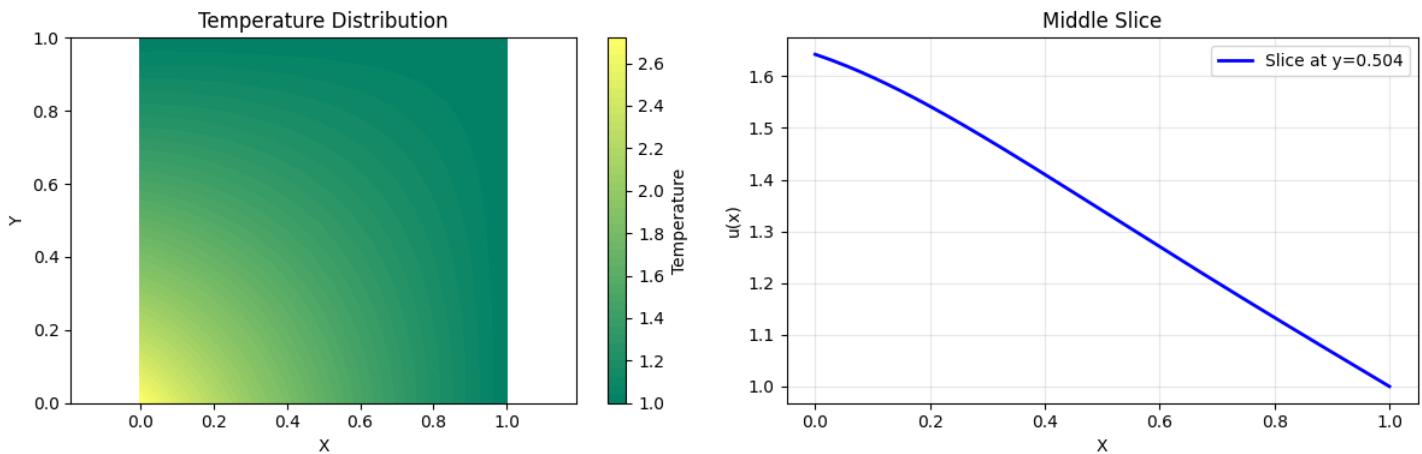
plt.tight_layout()
plt.show()

print(f"Min temperature: {u_grid.min():.4f}")
print(f"Max temperature: {u_grid.max():.4f}")
print(f"Center temperature: {u_grid[nx//2, ny//2]:.4f}")

except Exception as e:
    print(f"Error: {e}")

quick_plot()

```



Min temperature: 1.0000
Max temperature: 2.7183
Center temperature: 1.3374

In []:

gpu

In [14]:

```

# @title
%%file /tmp/t2v1.cu

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda_runtime.h>
#include <thrust/device_ptr.h>
#include <thrust/reduce.h>
#include <thrust/functional.h>
#include <algorithm>
#include <chrono>

// Параметры задачи (как в СРУ версии)
#define NX 128           // Количество узлов по x
#define NY 128           // Количество узлов по y
#define DX (1.0/(NX-1)) // Шаг по пространству
#define DY (1.0/(NY-1))
#define MAX_ITER 50000   // Максимальное количество итераций
#define TOLERANCE 1e-5    // Точность решения

```

```

// Размер блока
#define BLOCK_SIZE 16

#define CHECK_CUDA_ERROR(err) \
    if (err != cudaSuccess) { \
        printf("CUDA error: %s at %s:%d\n", cudaGetErrorString(err), __FILE__, __LINE__)
; \
        exit(1); \
    }

// CUDA kernel для одного шага Якоби
__global__ void jacobi_step_kernel(const double* u, double* u_new, int width, int height)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= 1 && i < width-1 && j >= 1 && j < height-1) {
        int idx = j * width + i;

        double dx2 = DX * DX;
        double dy2 = DY * DY;
        double factor = 0.5 / (1.0/dx2 + 1.0/dy2);

        double left = u[idx - 1];
        double right = u[idx + 1];
        double bottom = u[idx - width];
        double top = u[idx + width];

        u_new[idx] = factor * ((left + right) / dx2 + (bottom + top) / dy2);
    }
}

// CUDA kernel для вычисления невязки (норма С)
__global__ void compute_residual_kernel(const double* u, double* residual, int width, int height)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= 1 && i < width-1 && j >= 1 && j < height-1) {
        int idx = j * width + i;

        double dx2 = DX * DX;
        double dy2 = DY * DY;

        double laplacian = (u[idx - 1] - 2*u[idx] + u[idx + 1]) / dx2 +
                           (u[idx - width] - 2*u[idx] + u[idx + width]) / dy2;

        residual[idx] = fabs(laplacian);
    } else {
        // Для граничных точек невязка = 0
        int idx = j * width + i;
        residual[idx] = 0.0;
    }
}

// CUDA kernel для копирования граничных условий
__global__ void copy_boundaries_kernel(const double* u, double* u_new, int width, int height)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= width || j >= height) return;

    int idx = j * width + i;

    // Копируем граничные условия
    if (i == 0 || i == width-1 || j == 0 || j == height-1) {
        u_new[idx] = u[idx];
    }
}

```

```

// CUDA kernel для инициализации граничных условий
__global__ void initialize_boundaries_kernel(double* u, int width, int height) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if (i >= width || j >= height) return;

    int idx = j * width + i;

    // Нижняя граница: u(x,0) = e^(1-x)
    if (j == 0) {
        u[idx] = exp(1.0 - i * DX);
    }
    // Верхняя граница: u(x,1) = 1.0
    else if (j == height-1) {
        u[idx] = 1.0;
    }
    // Левая граница: u(0,y) = e^(1-y)
    else if (i == 0) {
        u[idx] = exp(1.0 - j * DY);
    }
    // Правая граница: u(1,y) = 1.0
    else if (i == width-1) {
        u[idx] = 1.0;
    }
    // Внутренние точки
    else {
        u[idx] = 0.0;
    }
}

// Kernel для вычисления разностей между массивами
__global__ void compute_diff_kernel(const double* a, const double* b, double* diff, int size) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;

    if (idx < size) {
        diff[idx] = fabs(a[idx] - b[idx]);
    }
}

// Функция для вычисления нормы С (максимум) с использованием Thrust
double compute_norm_c(const double* d_data, int size) {
    thrust::device_ptr<const double> thrust_ptr(d_data);
    return thrust::reduce(thrust_ptr, thrust_ptr + size, 0.0, thrust::maximum<double>());
}

// Функция для вычисления нормы L1 с использованием Thrust
double compute_norm_l1(const double* d_data, int size) {
    thrust::device_ptr<const double> thrust_ptr(d_data);
    double sum = thrust::reduce(thrust_ptr, thrust_ptr + size, 0.0, thrust::plus<double>());
    return sum * DX * DY;
}

// Функция для вычисления максимального изменения между итерациями
double compute_max_change(const double* d_u, const double* d_u_new, int size) {
    // Выделяем память для разностей
    double* d_diff;
    CHECK_CUDA_ERROR(cudaMalloc(&d_diff, size * sizeof(double)));

    // Вычисляем разности
    int blockSize1D = 256;
    int gridSize1D = (size + blockSize1D - 1) / blockSize1D;
    compute_diff_kernel<<<gridSize1D, blockSize1D>>>(d_u, d_u_new, d_diff, size);
    CHECK_CUDA_ERROR(cudaDeviceSynchronize());

    // Находим максимум с помощью Thrust
    double max_diff = compute_norm_c(d_diff, size);
}

```

```

// Освобождаем память
cudaFree(d_diff);

return max_diff;
}

// Функция для вычисления невязки (норма С)
double compute_residual(const double* d_u, int width, int height) {
    // Выделяем память для невязки
    double* d_residual;
    CHECK_CUDA_ERROR(cudaMalloc(&d_residual, width * height * sizeof(double)));

    // Настройка размеров блоков и гридов
    dim3 blockSize(BLOCK_SIZE, BLOCK_SIZE);
    dim3 gridSize((width + BLOCK_SIZE - 1) / BLOCK_SIZE,
                  (height + BLOCK_SIZE - 1) / BLOCK_SIZE);

    // Вычисляем невязку
    compute_residual_kernel<<<gridSize, blockSize>>>(d_u, d_residual, width, height);
    CHECK_CUDA_ERROR(cudaDeviceSynchronize());

    // Находим максимальную невязку
    double max_residual = compute_norm_c(d_residual, width * height);

    // Освобождаем память
    cudaFree(d_residual);

    return max_residual;
}

int main() {
    printf("Solving Laplace equation with CUDA\n");
    printf("Grid: %dx%d, dx=%f, dy=%f\n", NX, NY, DX, DY);
    printf("Tolerance: %.2e, Max iterations: %d\n", TOLERANCE, MAX_ITER);
    printf("Block size: %dx%d\n", BLOCK_SIZE, BLOCK_SIZE);

    // Выделение памяти на GPU
    double *d_u, *d_u_new;
    CHECK_CUDA_ERROR(cudaMalloc(&d_u, NX * NY * sizeof(double)));
    CHECK_CUDA_ERROR(cudaMalloc(&d_u_new, NX * NY * sizeof(double)));

    // Выделение памяти на CPU для вывода
    double* h_u = (double*)malloc(NX * NY * sizeof(double));

    // Настройка размеров блоков и гридов
    dim3 blockSize(BLOCK_SIZE, BLOCK_SIZE);
    dim3 gridSize((NX + BLOCK_SIZE - 1) / BLOCK_SIZE,
                  (NY + BLOCK_SIZE - 1) / BLOCK_SIZE);

    printf("Grid size: %dx%d blocks\n", gridSize.x, gridSize.y);

    // Инициализация
    initialize_boundaries_kernel<<<gridSize, blockSize>>>(d_u, NX, NY);
    initialize_boundaries_kernel<<<gridSize, blockSize>>>(d_u_new, NX, NY);
    CHECK_CUDA_ERROR(cudaDeviceSynchronize());

    auto start_computation = std::chrono::high_resolution_clock::now();

    // Основной итерационный цикл
    int iter = 0;
    double residual = 0.0;
    double max_change_val = 0.0;

    printf("\nStarting iterations...\n");
    printf("Iter Residual (C-norm)      Max-Change      L1-Norm\n");
    printf("-----\n");

    for (iter = 1; iter <= MAX_ITER; iter++) {
        // Шаг Якоби
        jacobi_step_kernel<<<gridSize, blockSize>>>(d_u, d_u_new, NX, NY);
        CHECK_CUDA_ERROR(cudaDeviceSynchronize());

```

```

// Копируем граничные условия
copy_boundaries_kernel<<<gridSize, blockSize>>>(d_u, d_u_new, NX, NY);
CHECK_CUDA_ERROR(cudaDeviceSynchronize());

// Вычисляем невязку и изменения
residual = compute_residual(d_u_new, NX, NY);
max_change_val = compute_max_change(d_u, d_u_new, NX * NY);
double l1_norm = compute_norm_l1(d_u_new, NX * NY);

// Меняем массивы местами
std::swap(d_u, d_u_new);

// Вывод прогресса каждые 1000 итераций
if (iter % 1000 == 0) {
    printf("%5d %12.6e %12.6e %12.6e\n",
           iter, residual, max_change_val, l1_norm);
}

// Критерий остановки (как в CPU версии)
if (residual < TOLERANCE && max_change_val < TOLERANCE) {
    printf("%5d %12.6e %12.6e %12.6e\n",
           iter, residual, max_change_val, l1_norm);
    printf("Converged!\n");
    break;
}
}

auto end_computation = std::chrono::high_resolution_clock::now();

if (iter > MAX_ITER) {
    printf("Reached maximum iterations without convergence\n");
    printf("Current residual: %.6e (target: %.1e)\n", residual, TOLERANCE);
}

// Копируем результат обратно на CPU
CHECK_CUDA_ERROR(cudaMemcpy(h_u, d_u, NX * NY * sizeof(double), cudaMemcpyDeviceToHost));

// Вывод результатов
printf("\n==== Final Results ====\n");
printf("Iterations: %d\n", iter);
printf("Final residual (C-norm): %.6e\n", residual);
printf("Final max change: %.6e\n", max_change_val);
printf("Final L1 norm: %.6e\n", compute_norm_l1(d_u, NX * NY));

// Сохранение в файл
FILE* file = fopen("solution_laplace_cuda.txt", "w");
if (file) {
    fprintf(file, "# X Y U\n");
    for (int j = 0; j < NY; j++) {
        for (int i = 0; i < NX; i++) {
            int idx = j * NX + i;
            fprintf(file, "%6f %6f %6f\n", i*DX, j*DY, h_u[idx]);
        }
    }
    fclose(file);
    printf("Saved to solution_laplace_cuda.txt\n");
}

// Анализ решения
printf("\n==== Solution Analysis ====\n");
printf("Boundary conditions check:\n");
printf("u(0,0) = %.6f (should be e = %.6f)\n", h_u[0], exp(1.0));
printf("u(1,0) = %.6f (should be 1.0)\n", h_u[NY-1], 1.0);
printf("u(0,1) = %.6f (should be 1.0)\n", h_u[(NY-1)*NX], 1.0);
printf("u(1,1) = %.6f (should be 1.0)\n", h_u[NY*NX-1], 1.0);

printf("\nInternal points:\n");
printf("u(0.5, 0.5) = %.6f\n", h_u[(NY/2)*NX + NX/2]);
printf("u(0.25, 0.25) = %.6f\n", h_u[(NY/4)*NX + NX/4]);
printf("u(0.75, 0.75) = %.6f\n", h_u[(3*NY/4)*NX + 3*NX/4]);

```

```

auto computation_time = std::chrono::duration_cast<std::chrono::microseconds>(
    end_computation - start_computation);

printf("\n==== CUDA Performance Results ====\n");
printf("Total computation time: %.3f ms\n", computation_time.count() / 1000.0);
printf("Time per iteration: %.3f microseconds\n", computation_time.count() / (double)iter);
printf("Iterations per second: %.0f\n", iter / (computation_time.count() / 1000000.0
));

// Освобождение памяти
free(h_u);
cudaFree(d_u);
cudaFree(d_u_new);

return 0;
}

```

Overwriting /tmp/t2v1.cu

In [15]:

```

# @title
!nvcc -arch=sm_75 /tmp/t2v1.cu -o /tmp/t2v1
# run the executable with nvprof
!nvprof /tmp/t2v1

```

```
/tmp/t2v1.cu(289): warning #225-D: the format string ends before this argument
    printf("u(1,0) = %.6f (should be 1.0)\n", h_u[128 -1], 1.0);
                                         ^

```

Remark: The warnings can be suppressed with "-diag-suppress <warning-number>"

```
/tmp/t2v1.cu(290): warning #225-D: the format string ends before this argument
    printf("u(0,1) = %.6f (should be 1.0)\n", h_u[(128 -1)*128], 1.0);
                                         ^

```

```
/tmp/t2v1.cu(291): warning #225-D: the format string ends before this argument
    printf("u(1,1) = %.6f (should be 1.0)\n", h_u[128*128 -1], 1.0);
                                         ^

```

```

/tmp/t2v1.cu: In function 'int main()':
/tmp/t2v1.cu:289:8: warning: too many arguments for format []8;;https://gcc.gnu.org/onlin
edocs/gcc/Warning-Options.html#index-Wformat-extra-args-Wformat-extra-args]8;;
 289 |     printf("u(1,0) = %.6f (should be 1.0)\n", h_u[NY-1], 1.0);
 |     ^~~~~~
/tmp/t2v1.cu:290:8: warning: too many arguments for format []8;;https://gcc.gnu.org/onlin
edocs/gcc/Warning-Options.html#index-Wformat-extra-args-Wformat-extra-args]8;;
 290 |     printf("u(0,1) = %.6f (should be 1.0)\n", h_u[(NY-1)*NX], 1.0);
 |     ^~~~~~
/tmp/t2v1.cu:291:8: warning: too many arguments for format []8;;https://gcc.gnu.org/onlin
edocs/gcc/Warning-Options.html#index-Wformat-extra-args-Wformat-extra-args]8;;
 291 |     printf("u(1,1) = %.6f (should be 1.0)\n", h_u[NY*NX-1], 1.0);
 |     ^~~~~~

```

Solving Laplace equation with CUDA

Grid: 128x128, dx=0.007874, dy=0.007874

Tolerance: 1.00e-05, Max iterations: 50000

Block size: 16x16

==1387== NVPROF is profiling process 1387, command: /tmp/t2v1

Grid size: 8x8 blocks

Starting iterations...

Iter	Residual (C-norm)	Max-Change	L1-Norm
1000	4.825725e+01	7.485691e-04	6.747362e-01
2000	2.282065e+01	3.537575e-04	8.815974e-01
3000	1.647062e+01	2.553533e-04	1.015109e+00
4000	1.249145e+01	1.936543e-04	1.109617e+00
5000	9.287306e+00	1.439822e-04	1.178395e+00
6000	6.855223e+00	1.062780e-04	1.228867e+00
7000	5.049931e+00	7.829375e-05	1.265995e+00

8000	3.719094e+00	5.765806e-05	1.293328e+00
9000	2.738657e+00	4.245810e-05	1.313454e+00
10000	2.016613e+00	3.126406e-05	1.328275e+00
11000	1.484947e+00	2.302152e-05	1.339188e+00
12000	1.093468e+00	1.695284e-05	1.347225e+00
13000	8.052044e-01	1.248398e-05	1.353143e+00
14000	5.929389e-01	9.193138e-06	1.357501e+00
15000	4.366326e-01	6.769775e-06	1.360711e+00
16000	3.215317e-01	4.985224e-06	1.363074e+00
17000	2.367732e-01	3.671091e-06	1.364814e+00
18000	1.743580e-01	2.703370e-06	1.366096e+00
19000	1.283961e-01	1.990746e-06	1.367039e+00
20000	9.455003e-02	1.465974e-06	1.367734e+00
21000	6.962606e-02	1.079535e-06	1.368246e+00
22000	5.127221e-02	7.949636e-07	1.368623e+00
23000	3.775656e-02	5.854066e-07	1.368901e+00
24000	2.780371e-02	4.310901e-07	1.369105e+00
25000	2.047449e-02	3.174523e-07	1.369255e+00
26000	1.507730e-02	2.337701e-07	1.369366e+00
27000	1.110284e-02	1.721470e-07	1.369448e+00
28000	8.176065e-03	1.267680e-07	1.369508e+00
29000	6.020808e-03	9.335126e-08	1.369552e+00
30000	4.433688e-03	6.874334e-08	1.369585e+00
31000	3.264943e-03	5.062220e-08	1.369609e+00
32000	2.404285e-03	3.727790e-08	1.369626e+00
33000	1.770502e-03	2.745123e-08	1.369639e+00
34000	1.303788e-03	2.021493e-08	1.369649e+00
35000	9.601018e-04	1.488616e-08	1.369656e+00
36000	7.070135e-04	1.096208e-08	1.369661e+00
37000	5.206407e-04	8.072417e-09	1.369665e+00
38000	3.833969e-04	5.944482e-09	1.369668e+00
39000	2.823313e-04	4.377483e-09	1.369670e+00
40000	2.079071e-04	3.223553e-09	1.369672e+00
41000	1.531016e-04	2.373806e-09	1.369673e+00
42000	1.127432e-04	1.748057e-09	1.369674e+00
43000	8.302344e-05	1.287260e-09	1.369674e+00
44000	6.113799e-05	9.479306e-10	1.369675e+00
45000	4.502170e-05	6.980512e-10	1.369675e+00
46000	3.315375e-05	5.140413e-10	1.369675e+00
47000	2.441420e-05	3.785368e-10	1.369675e+00
48000	1.797849e-05	2.787524e-10	1.369675e+00
49000	1.323928e-05	2.052720e-10	1.369676e+00
49918	9.997030e-06	1.550018e-10	1.369676e+00

Converged!

==== Final Results ===

Iterations: 49918

Final residual (C-norm): 9.997030e-06

Final max change: 1.550018e-10

Final L1 norm: 1.369676e+00

Saved to solution_laplace_cuda.txt

==== Solution Analysis ===

Boundary conditions check:

u(0,0) = 2.718282 (should be e = 2.718282)

u(1,0) = 1.000000 (should be 1.0)

u(0,1) = 1.000000 (should be 1.0)

u(1,1) = 1.000000 (should be 1.0)

Internal points:

u(0.5, 0.5) = 1.337383

u(0.25, 0.25) = 1.815447

u(0.75, 0.75) = 1.080496

==== CUDA Performance Results ===

Total computation time: 18305.850 ms

Time per iteration: 366.718 microseconds

Iterations per second: 2727

==1387== Profiling application: /tmp/t2v1

==1387== Profiling result:

Type	Time (%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	19.66%	758.31ms	99836	7.5950us	7.4870us	7.7750us	void cub::CU

B_200400_750_NS::DeviceReduceKernel<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>>>::Policy600, thrust::THRUST_200400_750_NS::device_ptr<double const >, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>, double, cuda::std::__4::__identity>(unsigned int, cub::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>>>::Policy600*, double, cub::CUB_200400_750_NS::GridEvenShare<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>>>::Policy600*)>, thrust::THRUST_200400_750_NS::maximum<double>, double const)	18.09%	697.57ms	49918	13.974us	13.887us	25.759us	compute_individual_kernel(double const *, double*, int, int)
	16.11%	621.17ms	49918	12.443us	12.383us	24.160us	jacobi_step_kernel(double const *, double*, int, int)
	13.97%	538.57ms	99836	5.3940us	5.0870us	5.8880us	void cub::CUB_200400_750_NS::DeviceReduceSingleTileKernel<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>>>::Policy600, double*, double*, int, thrust::THRUST_200400_750_NS::maximum<double>, double, double, cuda::std::__4::__identity>(unsigned int, double, thrust::THRUST_200400_750_NS::maximum<double>, cub::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::maximum<double>>>::Policy600, double*, int)
	9.40%	362.51ms	49919	7.2620us	7.1350us	7.4870us	void cub::CUB_200400_750_NS::DeviceReduceKernel<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::plus<double>>>::Policy600, thrust::THRUST_200400_750_NS::device_ptr<double const >, unsigned int, thrust::THRUST_200400_750_NS::plus<double>, double, cuda::std::__4::__identity>(unsigned int, cub::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::plus<double>>>::Policy600*, double, cub::CUB_200400_750_NS::GridEvenShare<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::plus<double>>>::Policy600*>, thrust::THRUST_200400_750_NS::plus<double>, double const)
	6.95%	267.84ms	49919	5.3650us	4.8320us	5.7270us	void cub::CUB_200400_750_NS::DeviceReduceSingleTileKernel<cuda::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::plus<double>>>::Policy600, double*, double*, int, thrust::THRUST_200400_750_NS::plus<double>, double, double, cuda::std::__4::__identity>(unsigned int, double, thrust::THRUST_200400_750_NS::plus<double>, cub::CUB_200400_750_NS::DeviceReducePolicy<double, unsigned int, thrust::THRUST_200400_750_NS::plus<double>>>::Policy600, double*, int)
	6.48%	249.98ms	149756	1.6690us	1.6310us	11.584us	[CUDA memcpy DtoH]
	5.57%	214.65ms	49918	4.3000us	4.1920us	4.4480us	compute_dif_f_kernel(double const *, double const *, double*, int)
	3.78%	145.69ms	49918	2.9180us	2.8480us	2.9760us	copy_boundaries_kernel(double const *, double*, int, int)
	0.00%	26.336us	2	13.168us	12.960us	13.376us	initialize_boundaries_kernel(double*, int, int)
	27.16%	4.18863s	499184	8.3900us	3.3270us	4.5818ms	cudaLaunchKernel API calls:
	18.58%	2.86506s	199673	14.348us	2.8330us	4.3109ms	cudaDevicesSynchronize
	15.86%	2.44611s	149755	16.334us	10.926us	5.0397ms	cudaMemcpyA sync
	10.48%	1.61643s	299510	5.3960us	1.4470us	6.1797ms	cudaStreamSync synchronize
	7.68%	1.18419s	249593	4.7440us	1.5750us	4.0849ms	cudaFree
	7.67%	1.18349s	249593	4.7410us	1.8420us	89.263ms	cudaMalloc
	6.16%	950.48ms	4193143	226ns	80ns	4.0196ms	cudaGetLastError
Error	2.20%	338.55ms	599021	565ns	198ns	4.0192ms	cudaGetDevice
	2.07%	319.12ms	299510	1.0650us	374ns	4.0188ms	cudaOccupancyMaxActiveBlocksPerMultiprocessorWithFlags
	1.34%	206.58ms	299510	689ns	216ns	4.0149ms	cudaDeviceGetAttribute
	0.81%	124.67ms	599020	208ns	82ns	3.5584ms	cudaPeekAtLastError
Attribute	0.00%	154.97us	114	1.3590us	118ns	67.864us	cuDeviceGetAttribute
	0.00%	131.17us	1	131.17us	131.17us	131.17us	cudaMemcpyName
	0.00%	13.715us	1	13.715us	13.715us	13.715us	cuDeviceGetName
Attributes	0.00%	8.8240us	1	8.8240us	8.8240us	8.8240us	cudaFuncGetAttributes
PCIBusId	0.00%	6.0470us	1	6.0470us	6.0470us	6.0470us	cuDeviceGetPCIBusId

	0.00%	1.7330us	3	577ns	159ns	1.3070us	cuDeviceGet
Count	0.00%	746ns	2	373ns	216ns	530ns	cuDeviceGet
	0.00%	641ns	1	641ns	641ns	641ns	cuDeviceTot
alMem	0.00%	462ns	1	462ns	462ns	462ns	cuDeviceGet
Uuid	0.00%	420ns	1	420ns	420ns	420ns	cuModuleGet
LoadingMode	0.00%	299ns	1	299ns	299ns	299ns	cudaGetDevic
ceCount	0.00%						

In [16]:

```
# @title
import numpy as np
import matplotlib.pyplot as plt

def quick_plot1():
    try:
        data = np.loadtxt('solution_laplace_cuda.txt', comments='#')
        x = data[:, 0]
        y = data[:, 1]
        u = data[:, 2]

        # Преобразование в сетку
        x_unique = np.unique(x)
        y_unique = np.unique(y)
        nx = len(x_unique)
        ny = len(y_unique)

        u_grid = u.reshape(nx, ny)
        X, Y = np.meshgrid(x_unique, y_unique, indexing='ij')

        plt.figure(figsize=(12, 4))

        plt.subplot(1, 2, 1)
        #plt.contourf(X, Y, u_grid, levels=50, cmap='cividis')
        plt.contourf(X, Y, u_grid, levels=50, cmap='summer')
        #plt.contourf(X, Y, u_grid, levels=50, cmap='Reds')
        plt.colorbar(label='Temperature')
        plt.xlabel('X')
        plt.ylabel('Y')
        plt.title('Temperature Distribution')
        plt.axis('equal')

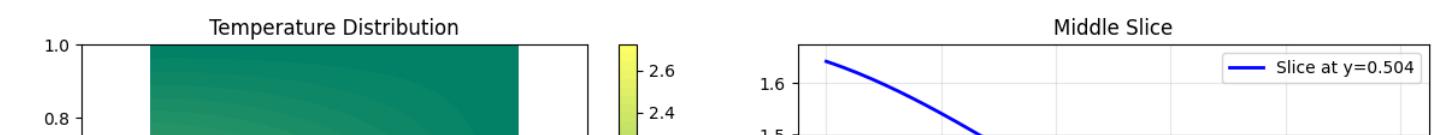
        plt.subplot(1, 2, 2)
        middle_y = ny // 2
        plt.plot(x_unique, u_grid[:, middle_y], 'b-', linewidth=2, label=f'Slice at y={y_unique[middle_y]:.3f}')
        plt.xlabel('X')
        plt.ylabel('u(x)')
        plt.title('Middle Slice')
        plt.grid(True, alpha=0.3)
        plt.legend()

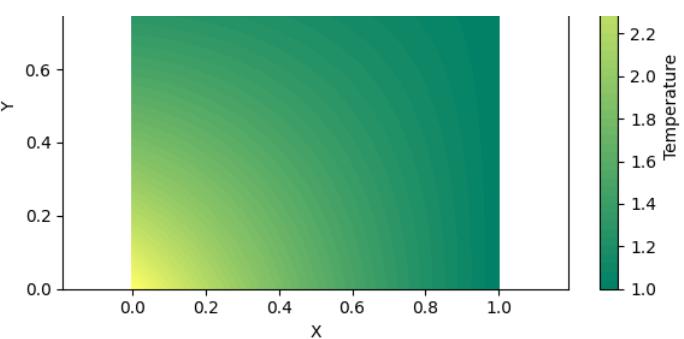
        plt.tight_layout()
        plt.show()

        print(f"Min temperature: {u_grid.min():.4f}")
        print(f"Max temperature: {u_grid.max():.4f}")
        print(f"Center temperature: {u_grid[nx//2, ny//2]:.4f}")

    except Exception as e:
        print(f"Error: {e}")

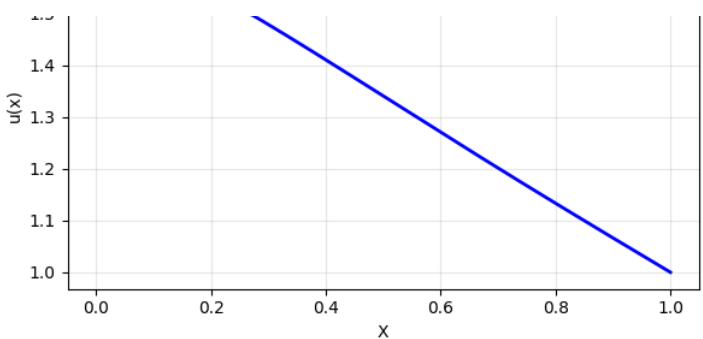
quick_plot1()
```





```
Min temperature: 1.0000  
Max temperature: 2.7183  
Center temperature: 1.3374
```

In []:



other

In []: